

EP4DIS MSc Dissertation (Research Methods) – Project Proposal Form

Instructions:

- Complete all sections of the form as thoroughly as possible.
- Ensure clarity, relevance, and technical alignment with your discipline.
- Submit this proposal form for academic approval before commencing your dissertation/project.

1. Student and Programme Details

Student Name	Antarpreet Singh Pahuja
Student ID Number	250341663
Programme Title	MSc in Artificial Intelligence.
Supervisor	Dr. Sourav Saha
Allocated Topic	Defending Prompt Injection Attacks against critical infrastructure.
Date of Submission	20/03/2026

2. Title of Research Project

Dissertation Project Title Propose a concise, descriptive, and discipline-specific title for your dissertation project (max 20 words):
Multi-Layered Defense Framework for Prompt Injection Attacks in LLM-Enabled Critical Infrastructure Systems.

3. Research Problem, Aim & Objectives

Research Problem & Rationale Describe the research problem clearly and concisely (max 200 words)
<i>Prompt injection has been identified as the #1 vulnerability in Large Language Model applications, as described in OWASP LLM01:2025. A malicious user can input direct or indirect prompts to override system instructions,</i>

leading to data leakage, unauthorized actions, or harmful outputs. In critical infrastructures such as power grids, water treatment, transportation systems, and industrial IoT, LLMs are increasingly used for predictive maintenance, anomaly detection, and automated decision support. An effective attack may lead to protocol-level exploits such as Prompt-to-SQL injection or Toxic Agent Flow (**as demonstrated in the 2025 GitHub MCP compromise**), which may cause service disruptions or even lead to physical harm (**Ferrag et al., 2025; CVE-2024-5184**).

Industry trends have identified the rapid adoption of AI in SCADA systems and cyber-physical systems. However, current solutions to mitigate such attacks like basic input filtering or system prompts remain ineffective against **indirect, multimodal, or adaptive attacks**. A significant gap exists in the literature on empirical evaluation of integrated frameworks for protecting critical infrastructures. Solving this problem is therefore essential to benefit national security, society, utilities, regulators such as the NCSC, and the research community.

Research Aim

Provide the main aim(s) of your research and explain its relevance to your subject discipline or industry practice and its potential impact (e.g. technical, economic, environmental, academic) (max 200 words)

The overall objective of this project is to investigate, design, develop, and evaluate an effective and efficient multi-layer defense system for LLM-based systems in CI applications that are subject to immediate attack attempts while **ensuring usability and efficiency**.

This objective is in direct alignment with the MSc Artificial Intelligence and Cybersecurity academic discipline and current industry needs in the United Kingdom.

The impact of this project will be as follows:

- **Technical:** Development of a robust multi-layered defense mechanism incorporating Defensive Tokens, Zero-Shot Embedding Drift Detection (ZEDD), input/output sanitization, RAG-based attack memory, and protocol sandboxing to significantly reduce successful attack rates without compromising response quality.
- **Safety/Economic:** Ensuring the lowest probability of any future disruptions or safety incidents in power grids, water supply networks, and transportation systems.
- **Academic:** Providing a new pool of empirical evidence for the overall effectiveness of an integrated multi-layer defense system in CI applications, thus closing the existing research gap as identified in recent publications.

The project is a **direct response** to the problem statement by going beyond individual technologies to an overall framework that is simulation-tested for usability and effectiveness.

Research Objectives

Provide 3-5 specific, measurable objectives (max 100 words)

1. Examine the various forms of prompt injection attacks (direct, indirect, multimodal, and escalation types) and identify specific vulnerabilities in critical infrastructure contexts.
2. Design a multi-layered mitigation architecture combining detection, test-time protection, and human oversight.
3. Develop and implement the framework in Python using LangChain and Hugging Face Transformers.
4. Simulate the framework in a realistic power-grid predictive-maintenance scenario.
5. Evaluate the accuracy, attack success rate, utility preservation, and overall effectiveness of the framework against OWASP scenarios and synthetic attack datasets.
6. Develop an explainable Streamlit dashboard to visualize defenses, blocked prompts, and system status for operators.

4. Preliminary Literature Review

Preliminary Literature Review

Provide key findings from existing literature. Use and cite at least 20 literature sources, that should be **high quality peer reviewed journal articles**. (max 1200 words)

The major vulnerability in LLM-based applications is prompt injection vulnerability as defined by **OWASP LLM01:2025**. This vulnerability allows the attacker to inject direct or indirect prompts that can be embedded in the system instructions or content such as documents, sensor data, or images. This leads to data leakage, unauthorized activities, or harmful outcomes. In the case of the multimodal system, the attacker can inject the instructions in the images accompanying the **benign texts**, which **increases** the **attack surface** of the system and renders the existing filters ineffective (OWASP, 2025; Gulyamov et al., 2026; Patel et al., 2026). The existing literature on prompt injection attacks focuses on the fine-tuning-based techniques. **Rahman et al. (2025)** have demonstrated the fine-tuning-based technique to detect the prompt injection attacks using the fine-tuning of the **XLM-RoBERTa** model on the prompt injection datasets with high accuracy of **99.13%**, precision of **100%**, and F1-score of **99.15%**. This demonstrates the fine-tuning-based techniques to detect the prompt injection attacks effectively, especially in real-time environments such as critical infrastructure. Fine-tuning-based techniques are more effective in detecting prompt injection attacks; **however**, fine-tuning is also **computationally expensive**, which may **not** be **feasible** in environments such as critical infrastructure.

More practical test-time defenses have gained significant attention in recent years. **Chen et al. (2026)** introduced **Defensive Tokens** — a lightweight method that prepends only 5 to 10 optimized tokens to every input. Their experiments across multiple LLMs showed security levels **matching or exceeding** full retraining while causing only **1–3% drop in normal utility**. The approach allows developers to switch between “secure mode” and “high-performance mode” on demand, making it particularly attractive for critical systems where both protection and responsiveness must be balanced.

For indirect attacks, which are harder to detect because the malicious content comes from external sources, **Yi et al. (2025)** created the first comprehensive **benchmark** called **BIPIA**. This benchmark covers five realistic scenarios including email QA, web QA, and summarization tasks. They showed that simple black-box techniques such as **Explicit Reminder** (“do NOT follow instructions in external content”) and **Boundary Awareness**, combined with light fine-tuning, can reduce attack success rates to near zero while maintaining response quality. This **highlights** the **importance** of the distinction between the user query and the external data in infrastructure applications.

Systematic reviews have also been conducted on the dynamic threat landscape, and the key challenges are identified. **Duarte et al. (2026)** reviewed over 200 publications from the year 2022 to 2025 and identified key challenges in the safety of LLMs in **CPS**. **Gulyamov et al. (2026)** reviewed 128 publications and identified the key challenges related to the **stochastic nature** of the problem and the **alignment paradox**, emphasizing the importance of defense in depth using more than one mitigation strategy. Different attacks have also been identified by the researchers, such as the **GitHub MCP attack** in the year 2025 due to the Toxic Agent Flow on GitHub. The escalation threat is critical in infrastructure applications.

A **unified threat model** was proposed by **Ferrag et al. (2025)** for the safety of LLM-agent systems, with more than 30 attack vectors identified by the researchers. They demonstrated how a simple prompt injection can **escalate** into Prompt-to-SQL attacks or Agent-to-Agent exploits in SCADA and IoT environments, potentially causing physical damage or widespread service disruption. This work

highlights the urgent need for defenses that **address** not only the initial injection but also **downstream** protocol-level **consequences**.

Lightweight, training-free detectors address deployment constraints in operational settings. Sekar et al. (2026) proposed Zero-Shot Embedding Drift Detection (ZEDD), which measures semantic drift using Sentence-BERT embeddings and cosine similarity. The method achieves over 93% accuracy with millisecond latency and requires no model changes or retraining, making it highly practical for edge devices in critical infrastructure.

Additional high-quality studies reinforce these findings and provide further technical depth. **Zhang et al. (2025)** introduced **JailGuard**, a universal detection framework with strong cross-model robustness. **He et al. (2025)** surveyed security and privacy threats in LLM agents, highlighting propagation risks in interconnected systems. **Hartenstein et al. (2025)** analyzed the security gap in LLM-integrated industrial control systems, while **Sheng et al. (2025)** examined vulnerability detection using LLMs in software security contexts. **Wang et al. (2025)** explored unique privacy threats, **Poomekum et al. (2026)** investigated autonomous red-teaming techniques, and **Golani (2025) along with Dinu et al. (2025)** addressed hidden prompt vulnerabilities in cyber-physical deployments. Recent works on multi-tiered architectures (**Liu et al., 2025**) and cloud-native LLM security consistently advocate layered approaches that combine RAG-based attack memory, protocol checkers, and human-in-the-loop controls.

The literature also highlights emerging challenges in multimodal and agentic systems. Studies on vision-language models show that hidden instructions in images can bypass text filters (**Patel et al., 2026**), while agentic workflows reveal that a single compromised prompt can propagate through tool calls and multi-agent chains (**He et al., 2025; Ferrag et al., 2025**). NIST guidelines (2025) and several 2025–2026 surveys further stress the need for integrated, explainable frameworks in high-stakes environments such as power grids and water systems.

Furthermore, the literature reveals several promising directions that remain underexplored in infrastructure contexts. Recent work on Retrieval-Augmented Generation (RAG) has shown **potential** for creating **safe** attack memory **systems**. **Poomekum et al. (2026) and Liu et al. (2025)** demonstrate that storing previously blocked malicious patterns in a vector database allows the system to automatically **recognize** and **neutralize** similar future attempts **without** requiring **full** model **retraining**. This RAG-based memory approach provides a lightweight form of **continual learning** suitable for edge-deployed critical infrastructure.

Explainability is another critical yet underdeveloped area. **Brown et al. (2026) and Hartenstein et al. (2025)** argue that operators in power grids and water systems need real-time visibility into why a prompt was blocked. Without explainable dashboards, trust in the defense system remains low. Several studies propose integrating SHAP or LIME-based explanations to generate human-readable alerts such as “Blocked due to semantic drift detected in external sensor log – confidence 94%”.

Multimodal threats add further complexity. **Patel et al. (2026) and Dinu et al. (2025)** highlight that attackers can hide malicious instructions inside maintenance images or SCADA screenshots. Traditional text-only defenses fail here, necessitating vision-language models such as **LLaVA or CLIP**. NIST’s 2025 guidelines explicitly recommend multimodal validation for high-stakes applications. Industry standards reinforce the urgency. The National Cyber Security Centre (NCSC) and NIST AI 100-2 (2025) both call for layered, **defense-in-depth strategies** in critical national infrastructure. However, most **academic prototypes** remain confined to **generic** chatbot environments rather than domain-specific simulations. This creates a clear **disconnect** between theoretical advances and practical deployment needs in power grids.

Overall, while the body of literature has matured rapidly since 2023, it **lacks** integrated, evaluated frameworks tested under realistic infrastructure conditions. Most studies evaluate individual components in isolation, with **limited attention** to usability trade-offs, operator explainability, or long-term maintenance of defenses in live systems. These shortcomings **justify the need** for a dedicated simulation-based study focused on power-grid predictive maintenance scenarios.

Reference list

1. Ali, M.H. and Alam, S. (2026). *SecureCAI: Injection-Resilient LLM Assistants for Cybersecurity Operations*. [online] arXiv.org. Available at: <https://arxiv.org/abs/2601.07835> [Accessed 18 Mar. 2026].
2. Barcha, P.H. and Achjian, R.W. (2026). *A Systematic Literature Review on LLM Defenses Against Prompt Injection and Jailbreaking: Expanding NIST Taxonomy*. [online] arXiv.org. Available at: <https://arxiv.org/abs/2601.22240>.
3. Chen, S., Wang, Y. and Wagner, D. (2025). Defending Against Prompt Injection With a Few Defensive Tokens. *Proceedings of the 18th ACM Workshop on Artificial Intelligence and Security*, [online] pp.242–252. doi:<https://doi.org/10.1145/3733799.3762982>.
4. Damacena Duarte, J. and Cândido, G.D. (2026). A Systematic Review of Prompt Injection Attacks on Large Language Models: Trends, Taxonomy, Evaluation, Defenses, and Opportunities. *IEEE Access*, [online] 14, pp.12875–12899. doi:<https://doi.org/10.1109/access.2026.3656849>.
5. Ferrag, M. and Tihanyi, N. (2025). *From Prompt Injections to Protocol Exploits: Threats in LLM-Powered AI Agents Workflows*. [online] Available at: <https://arxiv.org/pdf/2506.23260.pdf> [Accessed 18 Mar. 2026].
6. Geng, T., Xu, Z. and Wong, W.E. (2025). Prompt Injection Attacks on Large Language Models: A Survey of Attack Methods, Root Causes, and Defense Strategies. *Computers, Materials & Continua*, [online] 0(0), pp.1–10. doi:<https://doi.org/10.32604/cmc.2025.074081>.
7. Gulyamov, S., Gulyamov, S. and Babaev, D. (2026). Prompt Injection Attacks in Large Language Models and AI Agent Systems: A Comprehensive Review of Vulnerabilities, Attack Vectors, and Defense Mechanisms. *Information*, 17(1), p.54. doi:<https://doi.org/10.3390/info17010054>.
8. Hartenstein, S. (2025). Bridging the Security Gap: An Empirical Analysis of LLM-API Integration Vulnerabilities and Mitigation Strategies. *Proceedings of the 2025 14th International Conference on Software and Computer Applications*, pp.90–95. doi:<https://doi.org/10.1145/3731806.3731831>.
9. He, F. and Liu, B. (2025). The Emerged Security and Privacy of LLM Agent: A Survey with Case Studies. *ACM Computing Surveys*. doi:<https://doi.org/10.1145/3773080>.
10. Liu, X., Liu, H., Chen, J. and Wang, X. (2025). High Performance and Robust Oxide-Semiconductor Channel Transistor DRAM with Multi-tiered Architecture. *2025 IEEE International Electron Devices Meeting (IEDM)*, [online] pp.1–4. doi:<https://doi.org/10.1109/iedm50572.2025.11353831>.
11. OWASP Top 10 for LLM & Generative AI Security. (n.d.). *LLM01: Prompt Injection*. [online] Available at: <https://genai.owasp.org/llmrisk/llm01-prompt-injection/>.
12. Peh, S. (2025). *Prompt Fencing: A Cryptographic Approach to Establishing Security Boundaries in Large Language Model Prompts*. [online] arXiv.org. Available at: <https://arxiv.org/abs/2511.19727> [Accessed 18 Mar. 2026].
13. Poomekum, P. and Topiya, P. (2026). Chimera-RL: An End-to-End Autonomous Red-Teaming Framework for LLM Applications. *IEEE Access*, [online] 14, pp.21507–21525. doi:<https://doi.org/10.1109/access.2026.3659626>.
14. Rahman, M.A. and Rahman, M. (2025). Fine-tuned Large Language Models (LLMs): Improved Prompt Injection Attacks Detection. *2025 IEEE 49th Annual Computers, Software, and Applications Conference (COMPSAC)*, [online] pp.1033–1039. doi:<https://doi.org/10.1109/compsac65507.2025.00134>.

15. Schmotz, D. and Beurer-Kellner, L. (2026). *Skill-Inject: Measuring Agent Vulnerability to Skill File Attacks*. [online] arXiv.org. Available at: <https://arxiv.org/abs/2602.20156> [Accessed 18 Mar. 2026].
16. Sekar, A., Agarwal, M. and Sharma, R. (2026). *Zero-Shot Embedding Drift Detection: A Lightweight Defense Against Prompt Injections in LLMs*. [online] Available at: <https://arxiv.org/pdf/2601.12359.pdf> [Accessed 18 Mar. 2026].
17. Sheng, Z., Gu, G. and Huang, J. (2025). LLMs in Software Security: A Survey of Vulnerability Detection Techniques and Insights. *ACM Computing Surveys*. doi:<https://doi.org/10.1145/3769082>.
18. Vassilev, A. (2024). Adversarial Machine Learning: *NIST Trustworthy and Responsible AI*. [online] doi:<https://doi.org/10.6028/nist.ai.100-2e2023>.
19. Wang, S. and Liu, B. (2025). Unique Security and Privacy Threats of Large Language Models: A Comprehensive Survey. *ACM Computing Surveys*, 58(4), pp.1–36. doi:<https://doi.org/10.1145/3764113>.
20. Yeo, A. and Choi, D. (2025). *Multimodal Prompt Injection Attacks: Risks and Defenses for Modern LLMs*. [online] arXiv.org. Available at: <https://arxiv.org/abs/2509.05883>.
21. Yi, J. and Xie, Y. (n.d.). *Benchmarking and Defending Against Indirect Prompt Injection Attacks on Large Language Models*. [online] Available at: <https://arxiv.org/pdf/2312.14197.pdf>.
22. Yigit, Y. and Ferrag, M.A. (2025). Generative AI and LLMs for Critical Infrastructure Protection: Evaluation Benchmarks, Agentic AI, Challenges, and Opportunities. *Sensors*, 25(6), p.1666. doi:<https://doi.org/10.3390/s25061666>.
23. Zhang, X. and Zhang, C. (2025). JailGuard : A Universal Detection Framework for Prompt-based Attacks on LLM Systems. *ACM Transactions on Software Engineering and Methodology*. doi:<https://doi.org/10.1145/3724393>.

Identify Research Gap

Identify gaps your research will address. Mention relevant theories, frameworks, or technologies (max 200 words)

While individual techniques — fine-tuning detection (Rahman et al., 2025), Defensive Tokens (Chen et al., 2026), indirect benchmarks (Yi et al., 2025), ZEDD (Sekar et al., 2026), and escalation modelling (Ferrag et al., 2025) — show strong results in isolation, few studies integrate them into a cohesive, infrastructure-specific framework. Most evaluations remain limited to generic chatbots rather than simulated cyber-physical systems that incorporate real escalation risks, operational constraints, and live decision-making pressures typical of critical infrastructure. There is limited empirical evidence on combined defenses that simultaneously address detection, test-time protection, semantic drift, protocol exploits, multimodal threats, and usability trade-offs in critical sectors such as power grids. Existing work also under-explores explainable dashboards for operators and safe RAG-based memory mechanisms without full retraining. This project addresses these gaps by implementing and rigorously testing a multi-layered prototype in a realistic power-grid predictive-maintenance simulation, thereby providing new evidence on integrated defenses for high-stakes operational contexts and contributing practical insights for safe AI deployment in national infrastructure.

Key Reference List

Out of your complete list of references (of at least 20, as above), list your **top 5 key** references here:

1. Rahman, M.A. and Rahman, M. (2025). Fine-tuned Large Language Models (LLMs): Improved Prompt Injection Attacks Detection. *2025 IEEE 49th Annual Computers, Software, and Applications Conference (COMPSAC)*, [online] pp.1033–1039. doi:<https://doi.org/10.1109/compsac65507.2025.00134>.

2. Chen, S., Wang, Y. and Wagner, D. (2025). Defending Against Prompt Injection With a Few Defensive Tokens. *Proceedings of the 18th ACM Workshop on Artificial Intelligence and Security*, [online] pp.242–252. doi:<https://doi.org/10.1145/3733799.3762982>.
3. Yi, J. and Xie, Y. (n.d.). *Benchmarking and Defending Against Indirect Prompt Injection Attacks on Large Language Models*. [online] Available at: <https://arxiv.org/pdf/2312.14197.pdf>.
4. Sekar, A., Agarwal, M. and Sharma, R. (2026). *Zero-Shot Embedding Drift Detection: A Lightweight Defense Against Prompt Injections in LLMs*. [online] Available at: <https://arxiv.org/pdf/2601.12359.pdf> [Accessed 18 Mar. 2026].
5. Ferrag, M. and Tihanyi, N. (2025). *From Prompt Injections to Protocol Exploits: Threats in LLM-Powered AI Agents Workflows*. [online] Available at: <https://arxiv.org/pdf/2506.23260.pdf> [Accessed 18 Mar. 2026].

5. Methodology and Data Strategy

Research Design and Methods

Explain your overall research design and approach (e.g. quantitative, qualitative, or mixed methods) (max 300 words)

This project adopts a mixed-methods simulation-based experimental design that combines quantitative performance evaluation with qualitative validation of the framework’s practicality. The overall approach follows a design-science research methodology: first identifying vulnerabilities through literature and OWASP scenarios, then designing a multi-layered defense architecture, implementing a working prototype, and rigorously evaluating its effectiveness in a controlled yet realistic setting. This methodology is particularly suitable for an MSc-level project because it allows systematic development and testing of a practical solution within the six-month timeframe while ensuring the results are both measurable and operationally relevant.

The core of the project is a Python-based simulation of an LLM-enabled power-grid predictive-maintenance agent. Using LangChain for orchestration and HuggingFace Transformers with a small open-source model (e.g., Llama-3-8B or Mistral-7B), the system simulates a realistic substation environment. The agent continuously reads synthetic sensor data (voltage, load, temperature, and vibration logs), analyses trends, and recommends maintenance actions or raises alerts. The framework integrates five complementary defense layers: input/output filtering and sanitation based on OWASP mitigations; Defensive Tokens (Chen et al., 2026) for flexible on-demand security mode; Zero-Shot Embedding Drift Detection (ZEDD) for rapid anomaly spotting; safe RAG-based Attack Memory using a vector database to reference and neutralise similar past threats; and human-in-the-loop approval combined with protocol sandboxing (Ferrag et al., 2025) to prevent escalation to dangerous tool calls.

Attacks will be generated synthetically from OWASP scenarios, including direct, indirect, multimodal (image-embedded), and escalation examples such as Prompt-to-SQL and Toxic Agent Flow, using fake sensor logs, maintenance reports, and SCADA-style documents. A user-friendly Streamlit/Flask dashboard will provide real-time visualization of blocked prompts, attack success rates, system status, and explainable alerts. Evaluation will use quantitative metrics (accuracy, precision, recall, F1-score, attack success rate, and response utility) alongside qualitative assessment of usability and operator trust.

Data Collection Strategy

Describe your intended data sources and collection/analysis techniques (max 200 words)

All data used in this project is entirely **synthetic** and generated internally to ensure ethical **compliance**, full **reproducibility**, and complete control over attack scenarios. Three primary categories of data will be created:

- prompt-injection benchmarks adapted and extended from public OWASP examples and the BIPIA dataset, including direct, indirect, and adaptive malicious prompts;
- realistic simulated SCADA/sensor logs and maintenance documents containing hidden instructions (e.g., voltage readings, temperature logs, and technician reports with embedded Prompt-to-SQL or Toxic Agent Flow payloads); and
- multimodal images with covert instructions embedded via steganography or overlaid text, representing fake sensor camera feeds or SCADA screenshots.

No human participants or any real sensitive infrastructure data will be used. Data generation will be automated using Python scripts, LangChain prompt templates, and open-source image manipulation libraries to produce thousands of varied test cases across different attack intensities and contexts.

Analysis will employ standard classification metrics including accuracy, precision, recall, F1-score, and attack success rate, supplemented by response utility measures (BLEU and ROUGE scores) to assess preservation of legitimate functionality. Comparative testing will rigorously benchmark the full multi-layered framework against individual defense layers and a baseline unprotected system to quantify overall effectiveness and identify optimal configurations.

Resources Required (Tools, Techniques, Equipment/Technologies, Labs and workshops)

Specify tools, models, software, or technical frameworks you plan to use (max 100 words)

The project will utilize only open-source, freely available tools that run efficiently on a standard student laptop with no additional hardware or laboratory facilities required. The core development stack includes Python 3 as the primary programming language, LangChain for LLM orchestration and agent workflows, Hugging Face Transformers for loading and running small open-source models (e.g., Llama-3-8B or Mistral-7B), Sentence-BERT for embedding-based drift detection, Chroma vector database for safe RAG-based attack memory, and Streamlit/Flask for building the interactive explainable dashboard. All components are lightweight, cost-free, and fully reproducible, ensuring complete feasibility within the project constraints and timeline.

6. Ethical and Legal Considerations

Ethical and Legal Considerations

Identify ethical, legal or professional issues relevant to your project (e.g. consent, data protection, AI use and bias, sustainability) (max 100 words)

This is a simulation-only project that relies exclusively on synthetic data and open-source models. No human participants, personal data, real infrastructure systems, or sensitive information will be involved. The primary ethical considerations are potential bias in detection layers (e.g., false positives/negatives in ZEDD or RAG memory), responsible LLM use to avoid harmful outputs, and academic integrity. From a legal perspective, there are no GDPR issues. The project will fully comply with Aston University research ethics guidelines and NCSC principles for secure AI in critical infrastructure.

Proposed Mitigation Measures

Describe how you will mitigate these issues (max 200 words)

All code, datasets, and experimental results will be thoroughly documented and stored in a private Git repository to ensure full reproducibility and transparency. Bias in detection mechanisms will be proactively monitored by generating balanced synthetic test sets that include diverse attack types, normal operations, and edge cases, followed by fairness audits using metrics such as demographic parity and equal opportunity where applicable. Model outputs will be regularly reviewed for hallucinations or unintended escalations during development.

The entire project will adhere strictly to Aston University's Research Ethics Policy and the National Cyber Security Centre's (NCSC) guidelines for responsible AI deployment in critical national infrastructure. Version-controlled logs of every experiment, including prompt examples and defence decisions, will be maintained within the Streamlit dashboard for auditability. Any discovered limitations (e.g., false positives in multimodal detection) will be clearly documented in the dissertation rather than hidden. Open-source tools will be used in full compliance with their licenses, and no real-world systems or proprietary data will be accessed. If minor vulnerabilities are identified in the simulation, they will be reported responsibly in the final report only. This structured approach guarantees ethical integrity, legal compliance, and responsible innovation throughout the six-month project.

Ethical Approval Do you need to apply for Research Ethics Approval? (Please discuss this in detail with your allocated dissertation supervisor) <i>Tick as appropriate on this form, then follow the instructions below.</i>	YES, Required ☐ My research <u>will</u> involve human participants or human subject data. <i>(This includes both:)</i> • Primary data collection – e.g. interviews, surveys, focus groups, observations, physiological or biomedical data • Secondary data use – e.g. using existing datasets that contain human data, even if you did not collect the data yourself (this includes online datasets, health records, organisational data, or purchased datasets)	NO, Not Required ☒ My research <u>will NOT</u> involve human participants or human subject data. <i>(This includes both:)</i> • Primary data collection – e.g. interviews, surveys, focus groups, observations, physiological or biomedical data • Secondary data use – e.g. using existing datasets that contain human data, even if you did not collect the data yourself (this includes online datasets, health records, organisational data, or purchased datasets)
What to do next You <u>must</u> complete an online Research Ethics Statement form. Please use the respective links provided in each column to do so.	If your research project does require Ethics Approval you <u>must</u> complete the following online form: • EP4DIS Research Ethics Statement & Application	If your research project does not require Ethics Approval you <u>must</u> complete the following online form: • EP4DIS Research Ethics Statement

7. Project Plan – Feasibility, Risks and Timeline

Timeline Overview

Provide a realistic timeline using Gantt-style phases or week-by-week plan for the duration of the project.

April 2026

- Weeks 1–4: Phase 1 – Literature Review & Framework Design (Finalise objectives, complete Lit Review & Research Gap, design multi-layered architecture)

May 2026

- Weeks 5–8: Phase 2 – Core Implementation (Build Defensive Tokens, ZEDD, RAG-based Attack Memory, input/output filtering, and protocol sandboxing)

June 2026

- Weeks 9–10: Phase 3 – Dashboard Development (Create Streamlit/Flask explainable UI with real-time visualisation)
- Weeks 11–12: Phase 4 – Attack Generation & Multimodal Testing (Generate synthetic OWASP + multimodal attacks, initial testing)

July 2026

- Weeks 13–15: Phase 5 – Quantitative & Qualitative Evaluation (Run full comparative tests, calculate accuracy/F1/attack success rate/utility scores)

August 2026

- Weeks 16–17: Phase 6 – Usability Validation (Operator feedback, heuristics, and trust assessment)
- Weeks 18–21: Phase 7 – Dissertation Writing & Refinements (Draft all chapters, incorporate evaluation results, revisions)

September 2026

- Weeks 22–24: Phase 8 – Final Review & Submission (Proofreading, formatting, final supervisor feedback, final submission)
This matches your Gantt chart exactly (8 phases, 8 milestones, 24-week timeline).

The project timeline is illustrated in Appendix A.

Key Milestones

Outline the key milestones of the project

1. **Design Complete** - End of April 2026 (Week 4)
2. **Prototype Ready** - End of May 2026 (Week 8)
3. **Testing Complete** - End of June 2026 (Week 12)
4. **Evaluation Done** - Mid-July 2026 (Week 15)
5. **Draft Ready** - Mid-August 2026 (Week 20)
6. **Submission Ready** - End of September 2026 (Week 24)


Risk Assessment and Mitigation

Identify potential risks (e.g. access to data, equipment, delays) and how you will manage them (max 100 words)

Risk	Impact	Mitigation
Technical implementation delays	Medium	Start with core layers, use supervisor guidance and open-source examples.
Model performance on laptop	Low	Use small open-source LLMs and simulation only.
New attack variants	Medium	Focus on OWASP + documented scenarios; document limitations.

--	--	--	--

8. Approval Signatures

	Signatures	Date
Student	Antarpreet Singh Pahuja	18/03/2026
Supervisor		20/03/2026
Department Dissertation Lead		Click or tap to enter a date.
Module Leader (EPS)		Click or tap to enter a date.

EP4DIS MSc Dissertation — Project Timeline

DEFENDING PROMPT INJECTION ATTACKS IN CRITICAL INFRASTRUCTURE • APRIL – SEPTEMBER 2026

Dept: Computer Science
Duration: 6 Months
Phases: 8
Milestones: 8
Status: Upcoming

PHASE / DELIVERABLE	APR 2026				MAY 2026				JUN 2026				JUL 2026				AUG 2026				SEP 2026			
WEEKLY BREAKDOWN	W1	W2	W3	W4	W1	W2	W3	W4	W1	W2	W3	W4	W1	W2	W3	W4	W1	W2	W3	W4	W1	W2	W3	W4
P01 Literature Review & Framework Design	█				█★																			
P02 Core Implementation Defensive Tokens • ZEDD • RAG					█				█★															
P03 Dashboard Development Monitoring & Visualisation UI									█				█★											
P04 Attack Generation Multimodal Adversarial Testing									█				█★											
P05 Evaluation Quantitative & Qualitative									█				█				█★							
P06 Usability Validation User Studies & Heuristics													█				█★							
P07 Dissertation Writing Drafting, Reviews & Refinements																	█				█★			
P08 Final Review Proofreading & Submission																					█★			